

# Chapter 7 Moving Beyond Linearity

2023-02-05

```
library(splines)
library(ISLR2)
library(boot)
library(gam)
```

```
## Loading required package: foreach
```

```
## Loaded gam 1.22-1
```

```
library(ggplot2)
library(gridExtra)
library(splines)
library(leaps)
library(gam)
```

## Exercise 1

**a**

$$f_1(x) = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3$$

Here

$$a_1, b_1, c_1, d_1$$

are

$$\beta_0, \beta_1, \beta_2, \beta_3$$

respectively. And,

$$\beta_4 = 0$$

**b**

$$f(x) = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3 + \beta_4 (x^3 - 3x^2 \xi + 3x \xi^2 - \xi^3) = \beta_0 - \beta_4 \xi^3 + (\beta_1 + 3\xi^2 \beta_4)x + (\beta_2 - 3\xi \beta_4)x^2 + (\beta_3 + \beta_4)x^3$$

$$f_2(x) = \beta_0 - \beta_4 \xi^3 + (\beta_1 + 3\xi^2 \beta_4)x + (\beta_2 - 3\xi \beta_4)x^2 + (\beta_3 + \beta_4)x^3$$

Here

$$a_2, b_2, c_2, d_2$$

are

$$(\beta_0 - \beta_4 \xi^3), (\beta_1 + 3\xi^2 \beta_4), (\beta_2 - 3\xi \beta_4), (\beta_3 + \beta_4)$$

respectively.

**c**

$$f_1(\xi) = \beta_0 + \beta_1\xi + \beta_2\xi^2 + \beta_3\xi^3$$

$$f_2(\xi) = \beta_0 - \beta_4\xi^3 + (\beta_1 + 3\xi^2\beta_4)\xi + (\beta_2 - 3\xi\beta_4)\xi^2 + (\beta_3 + \beta_4)\xi^3 = \beta_0 + \beta_1\xi + \beta_2\xi^2 + \beta_3\xi^3 = f_1(\xi)$$

Therefore,

$$f(x)$$

is continuous at

$$\xi$$

**d**

$$f'_1(x) = \beta_1 + 2\beta_2x + 3\beta_3x^2$$

$$\Rightarrow f'_1(\xi) = \beta_1 + 2\beta_2\xi + 3\beta_3\xi^2$$

$$f'_2(x) = \beta_1 + 3\beta_4\xi^2 + 2(\beta_2 - 3\xi\beta_4)x + 3(\beta_3 + \beta_4)x^2$$

$$\Rightarrow f'_2(\xi) = \beta_1 + 3\xi^2\beta_4 + 2(\beta_2 - 3\xi\beta_4)\xi + 3(\beta_3 + \beta_4)\xi^2 = \beta_1 + 2\beta_2\xi + 3\beta_3\xi^2$$

Hence,

$$f'_1(\xi) = f'_2(\xi) = \beta_1 + 2\beta_2\xi + 3\beta_3\xi^2$$

In other words,

$$f'(x)$$

is continuous at

$$\xi$$

**e**

$$f''_1(x) = 2\beta_2 + 6\beta_3x$$

$$\Rightarrow f''_1(\xi) = 2\beta_2 + 6\beta_3\xi$$

$$f''_2(x) = 2(\beta_2 - 3\xi\beta_4) + 6(\beta_3 + \beta_4)x$$

$$\Rightarrow f''_2(\xi) = 2(\beta_2 - 3\xi\beta_4) + 6(\beta_3 + \beta_4)\xi = 2\beta_2 + 6\beta_3\xi = f''_1(\xi)$$

That is,

$$f''(x)$$

is continuous at

$$\xi$$

## Exercise 2

$$\hat{g} = \operatorname{argmin}_g \left( \sum_{i=1}^n (y_i - g(x_i))^2 + \lambda \int \left[ g^{(m)}(x) \right]^2 dx \right)$$

**a**

$$\lambda = \infty, m = 0$$

In order for

$$\hat{g}$$

to be minimised,

$$g(x)$$

must be zero otherwise the second term

$$\lambda \int \left[ g^{(m)}(x) \right]^2 dx$$

becomes very large. Therefore,

$$\hat{g}(x) = 0$$

**b**

$$\lambda = \infty, m = 1$$

In order for

$$\hat{g}$$

to be minimised,

$$g'(x)$$

must be zero otherwise the second term

$$\lambda \int \left[ g^{(1)}(x) \right]^2 dx$$

becomes very large.

$$g'(x) = 0 \Leftrightarrow g(x) = c$$

Here

$$c$$

is a constant number.

$$\hat{g}(x)$$

is a horizontal line.

**c**

$$\lambda = \infty, m = 2$$

Using the idea from (b),

$$g''(x) = 0 \Leftrightarrow g(x) = ax + b$$

In this case,

$$\hat{g}(x)$$

is a straight line.

**d**

$$\lambda = \infty, m = 3$$

$$g'''(x) = 0 \Leftrightarrow g(x) = ax^2 + bx + c$$

And

$$\hat{g}(x)$$

is a quadratic line in this scenario.

**e**

$$\lambda = 0, m = 3$$

Now,

$$\hat{g}$$

can be written as

$$\hat{g} = \underset{g}{\operatorname{argmin}} \left( \sum_{i=1}^n (y_i - g(x_i))^2 \right)$$

Just overfit the data as much as possible, till

$$y_i = g_i \quad \forall i$$

In other words,


$$\hat{g}$$

is now a curve that interpolates all observations.


## Exercise 3

```
set.seed(1)

X = seq(-2, 2, by = 0.1)
e = rnorm(41, 0, sd = 0.1)
Y = 1 + 1 * X - 2 * (X - 1)^2 * I(X >= 1) + e
plot(X, Y)
```




```
dat = data.frame(X, Y)
fit = lm(Y ~ bs(X, knots = 1), data = dat)
pred = predict(fit, newdata = list(X = X))
plot(X, Y)
lines(X, pred, lwd = 2, col = 'blue')
```




## Exercise 4

```
set.seed(1)

X = seq(-2, 6, by = 0.1)
b1 = I(X >= 0 & X <= 2) - (X - 1) * I(X >= 1 & X <= 2)
b2 = (X - 3) * I(X >= 3 & X <= 4) + I(X > 4 & X <= 5)
e = rnorm(81, 0, sd = 0.1)
Y = 1 + 1 * b1 + 3 * b2 + e
plot(X, Y)
```



```
dat = data.frame(X, Y)
fit = lm(Y ~ bs(X, knots = c(0, 1, 2, 3, 4, 5)), data = dat)
pred = predict(fit, newdata = list(X = X))
plot(X, Y)
lines(X, pred, lwd = 2, col = 'blue')
```



## Exercise 5

$$\hat{g}_1 = \operatorname{argmin}_g \left( \sum_{i=1}^n (y_i - g(x_i))^2 + \lambda \int [g^{(3)}(x)]^2 dx \right)$$

$$\hat{g}_2 = \operatorname{argmin}_g \left( \sum_{i=1}^n (y_i - g(x_i))^2 + \lambda \int [g^{(4)}(x)]^2 dx \right)$$

From exercise 2, we know that when

$$\lambda$$

approaches infinity, the two curve functions above can be minimised by setting

$$g^{(m)} = 0$$

$$g^{(3)} = 0 \Leftrightarrow g_3(x) = ax^2 + bx + c$$

$$g^{(4)} = 0 \Leftrightarrow g_4(x) = ax^4 + bx^3 + cx + d$$



**a**

$$\hat{g}_4(x)$$

has the smaller training RSS as it is more flexible.

**b**

We can't say for sure when it comes to the test RSS, it depends on the true relationship between X and Y.

**c**

For

$$\lambda = 0$$

, the both two curves interpolate all its observations therefore there they will have the same training RSS which is zero. As for the test RSS, we can't say for sure also since it depends on the true relationship between X and Y.

## Exercise 6

```
head(Wage)
```

```
##      year age      maritl      race      education      region
## 231655 2006  18 1. Never Married 1. White      1. < HS Grad 2. Middle Atlantic
## 86582  2004  24 1. Never Married 1. White 4. College Grad 2. Middle Atlantic
## 161300 2003  45      2. Married 1. White 3. Some College 2. Middle Atlantic
## 155159 2003  43      2. Married 3. Asian 4. College Grad 2. Middle Atlantic
## 11443  2005  50      4. Divorced 1. White      2. HS Grad 2. Middle Atlantic
## 376662 2008  54      2. Married 1. White 4. College Grad 2. Middle Atlantic
##      jobclass      health health_ins logwage      wage
## 231655 1. Industrial      1. <=Good      2. No 4.318063 75.04315
## 86582  2. Information 2. >=Very Good      2. No 4.255273 70.47602
## 161300 1. Industrial      1. <=Good      1. Yes 4.875061 130.98218
## 155159 2. Information 2. >=Very Good      1. Yes 5.041393 154.68529
## 11443  2. Information      1. <=Good      1. Yes 4.318063 75.04315
## 376662 2. Information 2. >=Very Good      1. Yes 4.845098 127.11574
```

```
colnames(Wage)
```

```
## [1] "year"      "age"      "maritl"    "race"      "education"
## [6] "region"    "jobclass" "health"    "health_ins" "logwage"
## [11] "wage"
```


**a**

The optimal degree  $d = 9$  was chosen the most often. However, using anova, models with 4 degrees or 9 degrees were chosen. The results from the two approaches somewhat agree with each other.

CV

```
deltas = rep(NA, 10)
for (i in 1:10){
  model = glm(wage ~ poly(age, i), data = Wage)
  cv_model = cv.glm(model, data = Wage, K = 10)
  deltas[i] = cv_model$delta[1]
}

plot(1:10, deltas, type = 'l')
points(which.min(deltas), deltas[which.min(deltas)], col = 'red')
```



```
optimal_degrees = rep(NA, 3)

for (i in 1:100){

  deltas = rep(NA, 10)

  for (j in 1:10){

    model = glm(wage ~ poly(age, j), data = Wage)
    cv_model = cv.glm(model, data = Wage, K = 10)
    deltas[j] = cv_model$delta[1]
  }

  optimal_degrees[i] = which.min(deltas)
```

```

    optimal_degrees[i] = which.min(deltas)
}

```

```

table(optimal_degrees)

```

```

## optimal_degrees
##  4  5  6  7  8  9 10
##  6  4 10 17  7 35 21


```

```

agelims = range(Wage$age)
age_grid = seq(from = agelims[1], to = agelims[2])
preds = predict(glm(wage ~ poly(age, 9), data = Wage),
                 newdata = list(age = age_grid), se = T)
se_bands = cbind(preds$fit + 2 * preds$se.fit, preds$fit - 2 * preds$se.fit)

plot(Wage$age, Wage$wage, xlim = agelims, cex = .5, col = 'darkgrey')
lines(age_grid, preds$fit, lwd = 2, col = 'blue')
matlines(age_grid, se_bands, lwd = 1, col = 'blue', lty = 3)

```



ANOVA

```

fit1 = lm(wage ~ age, data = Wage)
fit2 = lm(wage ~ poly(age, 2), data = Wage)
fit3 = lm(wage ~ poly(age, 3), data = Wage)
fit4 = lm(wage ~ poly(age, 4), data = Wage)
fit5 = lm(wage ~ poly(age, 5), data = Wage)
fit6 = lm(wage ~ poly(age, 6), data = Wage)
fit7 = lm(wage ~ poly(age, 7), data = Wage)
fit8 = lm(wage ~ poly(age, 8), data = Wage)
fit9 = lm(wage ~ poly(age, 9), data = Wage)
fit10 = lm(wage ~ poly(age, 10), data = Wage)

anova(fit1, fit2, fit3, fit4, fit5, fit6, fit7, fit8, fit9, fit10)

```

```
## Analysis of Variance Table
```

```
##
```

```
## Model 1: wage ~ age
```

```
## Model 2: wage ~ poly(age, 2)
```

```
## Model 3: wage ~ poly(age, 3)
```

```
## Model 4: wage ~ poly(age, 4)
```

```
## Model 5: wage ~ poly(age, 5)
```

```
## Model 6: wage ~ poly(age, 6)
```

```
## Model 7: wage ~ poly(age, 7)
```

```
## Model 8: wage ~ poly(age, 8)
```

```
## Model 9: wage ~ poly(age, 9)
```

```
## Model 10: wage ~ poly(age, 10)
```

##	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
## 1	2998	5022216				
## 2	2997	4793430	1	228786	143.7638	< 2.2e-16 ***
## 3	2996	4777674	1	15756	9.9005	0.001669 **
## 4	2995	4771604	1	6070	3.8143	0.050909 .
## 5	2994	4770322	1	1283	0.8059	0.369398
## 6	2993	4766389	1	3932	2.4709	0.116074
## 7	2992	4763834	1	2555	1.6057	0.205199
## 8	2991	4763707	1	127	0.0796	0.777865
## 9	2990	4756703	1	7004	4.4014	0.035994 *
## 10	2989	4756701	1	3	0.0017	0.967529

```
## ---
```


```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```

preds = predict(glm(wage ~ poly(age, 3), data = Wage),
                newdata = list(age = age_grid), se = T)
se_bands = cbind(preds$fit + 2 * preds$se.fit, preds$fit - 2 * preds$se.fit)

plot(Wage$age, Wage$wage, xlim = agelims, cex = .5, col = 'darkgrey')
lines(age_grid, preds$fit, lwd = 2, col = 'blue')
matlines(age_grid, se_bands, lwd = 1, col = 'blue', lty = 3)

```




b

```
deltas = rep(NA, 10)

for (i in 2:11){
  Wage$age_cut = cut(Wage$age, i)
  model = glm(wage ~ age_cut, data = Wage)
  cv_model = cv.glm(Wage, model, K = 10)
  deltas[i-1] = cv_model$delta[1]
}

plot(2:11, deltas, type = 'l')
min_index = which.min(deltas)
points(min_index + 1, deltas[min_index], col = 'red')
```



```

optimal_cuts = rep(NA, 10)

for (i in 1:100){

  set.seed(i)
  deltas = rep(NA, 10)

  for (j in 2:11){
    Wage$age_cut = cut(Wage$age, j)
    model = glm(wage ~ age_cut, data = Wage)
    cv_model = cv.glm(Wage, model, K = 10)
    deltas[j-1] = cv_model$delta[1]
  }

  optimal_cuts[i] = which.min(deltas) + 1

}

table(optimal_cuts)

```

```

## optimal_cuts
## 8 11
## 22 78

```

## Exercise 7

### Exploring

Here I use four predictors, they are marital status, education, race and health. In terms of





Marital status: doesn't play have a strong relationship with wage. Married workers tend to have higher income on average. Most people whose income really high (higher than \$200,000) are in the never married and married group.

Education: on average, the higher in education level, the higher in wage.

Race: white and Asian people tend to have higher income on average. Most people whose income really high (higher than \$200,000) are in the white group.

Health: seems to be not a significant predictor. People with good health condition tend to have higher income on average.

```
par(mfrow = c(2, 2))
plot(Wage$maritl, Wage$wage)
plot(Wage$education, Wage$wage)
plot(Wage$race, Wage$wage)
plot(Wage$health, Wage$wage)
```




## Modelling

```
library(gam)
model_1 = gam(wage ~ lo(age, span = 0.7) + education, data = Wage)
model_2 = gam(wage ~ lo(age, span = 0.7) + education + maritl, data = Wage)
model_3 = gam(wage ~ lo(age, span = 0.7) + education + maritl + race, data = Wage)
model_4 = gam(wage ~ lo(age, span = 0.7) + education + maritl + race + health, data = Wage)


anova(model_1, model_2, model_3, model_4)
```


```
## Analysis of Deviance Table
##
## Model 1: wage ~ lo(age, span = 0.7) + education
## Model 2: wage ~ lo(age, span = 0.7) + education + maritl
## Model 3: wage ~ lo(age, span = 0.7) + education + maritl + race
## Model 4: wage ~ lo(age, span = 0.7) + education + maritl + race + health
##   Resid. Df Resid. Dev Df Deviance Pr(>Chi)
## 1      2992.8    3737372
## 2      2988.8    3634221  4   103150 < 2.2e-16 ***
## 3      2985.8    3626320  3     7901  0.08732 .
## 4      2984.8    3594978  1     31342 3.375e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```


```
plot(model_4)
```

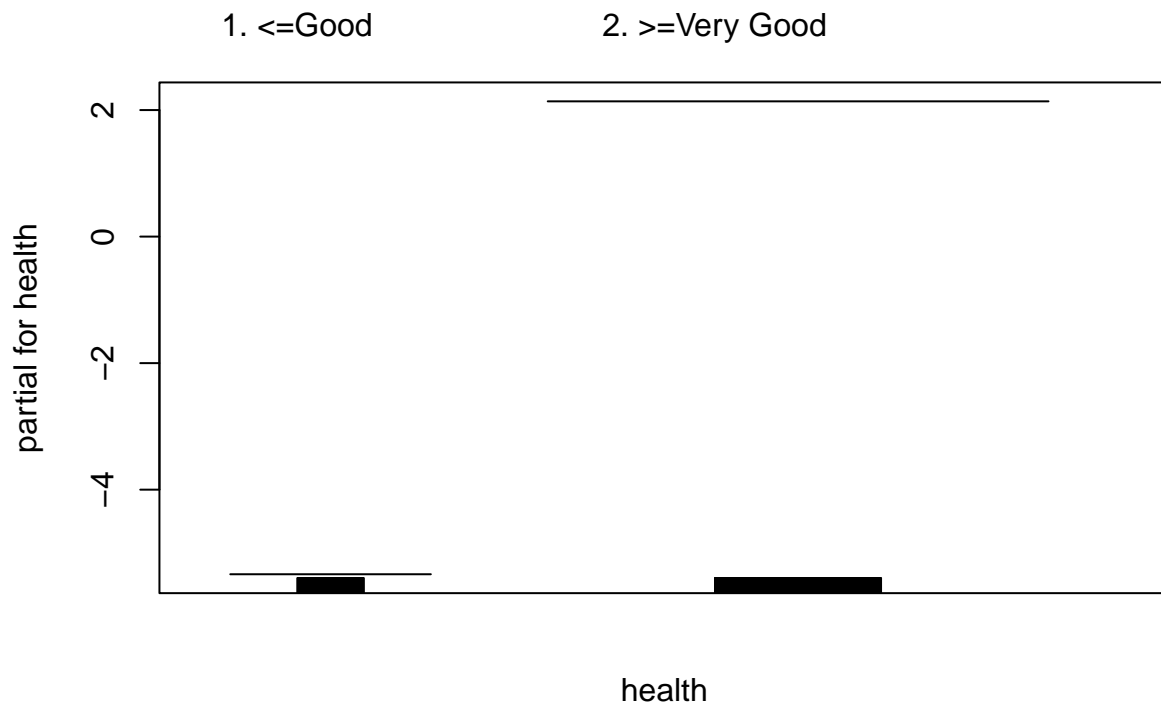












## Exercise 8

```
head(Auto)
```

```
##   mpg cylinders displacement horsepower weight acceleration year origin
## 1  18         8         307         130   3504          12.0    70      1
## 2  15         8         350         165   3693          11.5    70      1
## 3  18         8         318         150   3436          11.0    70      1
## 4  16         8         304         150   3433          12.0    70      1
## 5  17         8         302         140   3449          10.5    70      1
## 6  15         8         429         198   4341          10.0    70      1
##                                name
## 1 chevrolet chevelle malibu
## 2      buick skylark 320
## 3    plymouth satellite
## 4      amc rebel sst
## 5      ford torino
## 6    ford galaxie 500
```

```
colnames(Auto)
```



```
## [1] "mpg"          "cylinders"    "displacement" "horsepower"   "weight"
## [6] "acceleration" "year"         "origin"       "name"
```

There is absolutely non-linear relationship between some predictors and the response (here is mpg as we did in the previous chapters). We can plot the relationship, perform polynomial model or step function to ensure there is indeed a non-linear relationship.

## Plots

```
plot1 = ggplot(Auto, aes(x = displacement, y = mpg)) +  
  geom_point(alpha = 0.5) +  
  geom_smooth()  
  
plot2 = ggplot(Auto, aes(x = displacement, y = mpg)) +  
  geom_point(alpha = 0.5) +  
  geom_smooth()  
  
grid.arrange(plot1, plot2, ncol = 2)
```

```
## 'geom_smooth()' using method = 'loess' and formula = 'y ~ x'  
## 'geom_smooth()' using method = 'loess' and formula = 'y ~ x'
```



## CV


Here we see that the optimal degree is 10, significantly better than 1.

```

deltas = rep(NA, 10)
for (i in 1:10){
  model = glm(mpg ~ poly(displacement, i), data = Auto)
  cv_model = cv.glm(model, data = Auto, K = 10)
  deltas[i] = cv_model$delta[1]
}

plot(1:10, deltas, type = 'l')
points(which.min(deltas), deltas[which.min(deltas)], col = 'red')

```



## Cuts

It is suggested to have 9 cuts. (If the relationship is linear, the number of cuts should be small).

```


deltas = rep(NA, 10)

for (i in 2:11){
  Auto$dis_cut = cut(Auto$displacement, i)
  model = glm(mpg ~ dis_cut, data = Auto)
  cv_model = cv.glm(Auto, model, K = 10)
  deltas[i-1] = cv_model$delta[1]
}

plot(2:11, deltas, type = 'l')

```

```
min_index = which.min(deltas)
points(min_index + 1, deltas[min_index], col = 'red')
```




## Exercise 9

```
head(Boston)
```

```
##      crim zn  indus chas   nox    rm  age   dis rad tax ptratio lstat medv
## 1 0.00632 18  2.31    0 0.538 6.575 65.2 4.0900   1 296   15.3  4.98 24.0
## 2 0.02731  0  7.07    0 0.469 6.421 78.9 4.9671   2 242   17.8  9.14 21.6
## 3 0.02729  0  7.07    0 0.469 7.185 61.1 4.9671   2 242   17.8  4.03 34.7
## 4 0.03237  0  2.18    0 0.458 6.998 45.8 6.0622   3 222   18.7  2.94 33.4
## 5 0.06905  0  2.18    0 0.458 7.147 54.2 6.0622   3 222   18.7  5.33 36.2
## 6 0.02985  0  2.18    0 0.458 6.430 58.7 6.0622   3 222   18.7  5.21 28.7
```

a

```
cubic_model = lm(nox ~ poly(dis, 3), data = Boston)
plot(Boston$dis, Boston$nox, col = 'darkgrey')
points(cubic_model$fitted.values, type = 'l', col = 'blue')
```

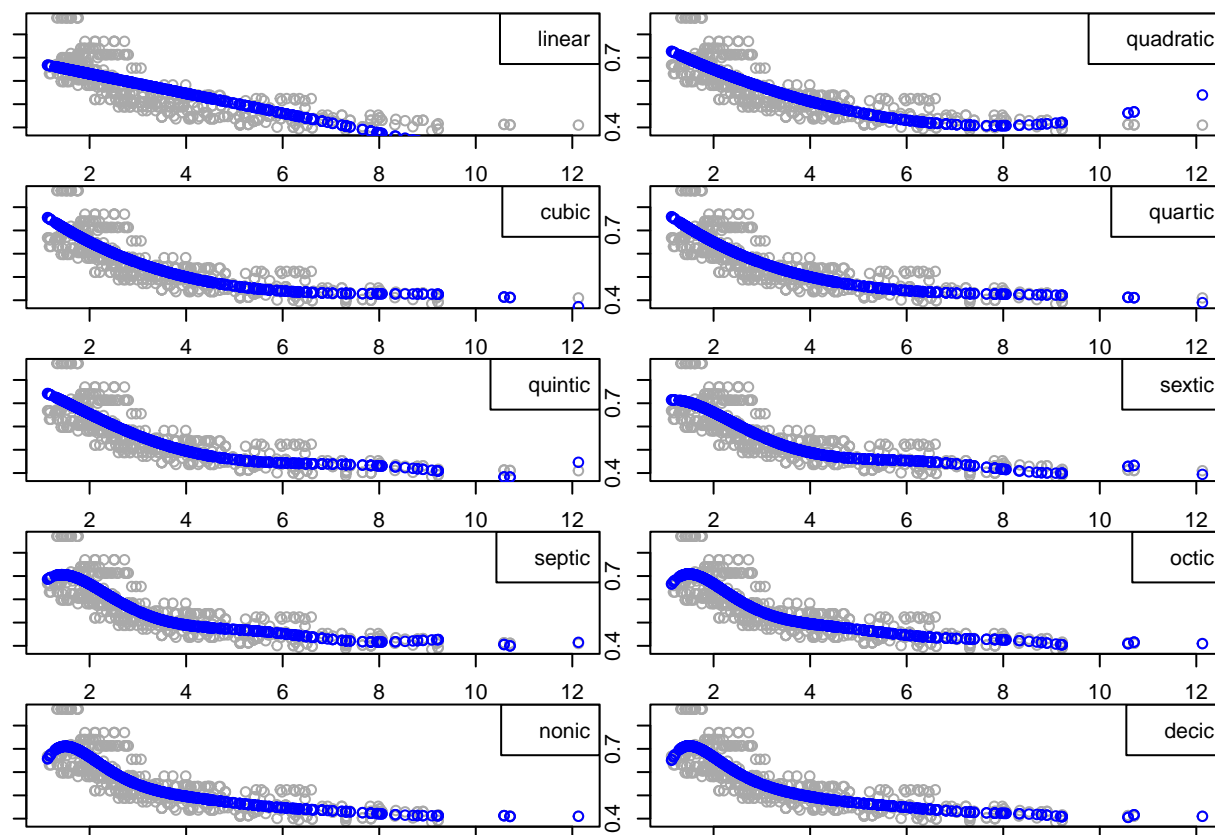


b


```
par(mfrow = c(5, 2))
par(mar = c(1, 1, 1, 1))
polynomials = c('linear', 'quadratic', 'cubic', 'quartic', 'quintic',
                 'sextic', 'septic', 'octic', 'nonic', 'decic')
RSS = rep(0, 10)

for (i in 1:10){
  model = lm(nox ~ poly(dis, i), data = Boston)
  plot(Boston$dis, Boston$nox, col = 'darkgrey')
  points(Boston$dis, model$fitted.values, col = 'blue')
  legend('topright', legend = polynomials[i])
  RSS[i] = sum(model$residuals^2)
}
```





```
plot(1:10, RSS, type = 'l')
legend('topright', legend = c('RSS'))
```




c

```
deltas = rep(0, 10)

for (i in 1:10){
  model = glm(nox ~ poly(dis, i), data = Boston)
  delta = cv.glm(Boston, model, K = 10)$delta[1]
  deltas[i] = delta
}

plot(1:10, deltas, type = 'l')
points(which.min(deltas), deltas[which.min(deltas)])
```



The optimal degree is four and this suggests that there is a non-linear relationship between the response the the predictor. From (b), we see that a decic model overfits the data. You might notice that a nonic model is also good enough but it is not the parsimonious model.

**d**


```
attr(bs(Boston$dis, df = 4), 'knots')
```

```
##      50%
```

```
## 3.20745
```

```
spline_model = lm(nox ~ bs(dis, df = 4), data = Boston)
spline_preds = predict(spline_model)
```

```
plot(Boston$dis, Boston$nox, col = 'darkgrey')
points(Boston$dis, spline_preds, col = 'blue')
```



e

Using the whole dataset and RSS as a metric, 10 degrees of freedom turns out to be the best number. This is predictable since our model tried to overfit the data.


```
RSS = rep(NA, 10)

for (i in 1:10){
  model = lm(nox ~ bs(dis, df = i), data = Boston)
  RSS[i] = sum(model$residuals^2)
}
```

```
## Warning in bs(dis, df = i): 'df' was too small; have used 3
```

```
## Warning in bs(dis, df = i): 'df' was too small; have used 3
```

```
plot(1:10, RSS, type = 'l')
```



f

CV

```
set.seed(1)
deltas = rep(NA, 16)

for (i in 3:18){
  model = glm(nox ~ bs(dis, df = i), data = Boston)
  cv_model = cv.glm(data = Boston, model, K = 10)
  deltas[i-2] = cv_model$delta[1]
}
```

```
## Warning in bs(dis, degree = 3L, knots = numeric(0), Boundary.knots = c(1.137, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases
```

```
## Warning in bs(dis, degree = 3L, knots = numeric(0), Boundary.knots = c(1.137, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases
```

```
## Warning in bs(dis, degree = 3L, knots = numeric(0), Boundary.knots = c(1.1296, :
```

```

## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = numeric(0), Boundary.knots = c(1.1296, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('50%' = 3.0993), Boundary.knots =
## c(1.137, : some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('50%' = 3.0993), Boundary.knots =
## c(1.137, : some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('50%' = 3.3603), Boundary.knots =
## c(1.1296, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases

## Warning in bs(dis, degree = 3L, knots = c('50%' = 3.3603), Boundary.knots =
## c(1.1296, : some 'x' values beyond boundary knots may cause ill-conditioned
## bases

## Warning in bs(dis, degree = 3L, knots = c('33.33333%' = 2.388766666666667, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('33.33333%' = 2.388766666666667, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('33.33333%' = 2.3088, '66.66667%'
## = 4.097266666666667: some 'x' values beyond boundary knots may cause
## ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('33.33333%' = 2.3088, '66.66667%'
## = 4.097266666666667: some 'x' values beyond boundary knots may cause
## ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('25%' = 2.087875, '50%' = 3.19095, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('25%' = 2.087875, '50%' = 3.19095, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('20%' = 1.92404, '40%' = 2.55946, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('20%' = 1.92404, '40%' = 2.55946, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('20%' = 1.94984, '40%' = 2.59774, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('20%' = 1.94984, '40%' = 2.59774, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

```

```

## Warning in bs(dis, degree = 3L, knots = c('16.66667%' = 1.866366666666667, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('16.66667%' = 1.866366666666667, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('16.66667%' = 1.82085, '33.33333%'
## = 2.363866666666667, : some 'x' values beyond boundary knots may cause
## ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('16.66667%' = 1.82085, '33.33333%'
## = 2.363866666666667, : some 'x' values beyond boundary knots may cause
## ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('14.28571%' = 1.79078571428571, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('14.28571%' = 1.79078571428571, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('14.28571%' = 1.7912, '28.57143%' =
## 2.1705, : some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('14.28571%' = 1.7912, '28.57143%' =
## 2.1705, : some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('12.5%' = 1.757275, '25%' = 2.1084, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('12.5%' = 1.757275, '25%' = 2.1084, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('12.5%' = 1.76375, '25%' = 2.10525, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('12.5%' = 1.76375, '25%' = 2.10525, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('11.11111%' = 1.691244444444444, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('11.11111%' = 1.691244444444444, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('11.11111%' = 1.712977777777778, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('11.11111%' = 1.712977777777778, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('10%' = 1.66236, '20%' = 1.98518, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

```

```

## Warning in bs(dis, degree = 3L, knots = c('10%' = 1.66236, '20%' = 1.98518, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('10%' = 1.6624, '20%' = 1.9769, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('10%' = 1.6624, '20%' = 1.9769, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('9.090909%' = 1.59007272727273, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('9.090909%' = 1.59007272727273, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('9.090909%' = 1.61941818181818, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('9.090909%' = 1.61941818181818, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('8.333333%' = 1.58948333333333, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('8.333333%' = 1.58948333333333, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('8.333333%' = 1.5874, '16.66667%' =
## 1.8651, : some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('8.333333%' = 1.5874, '16.66667%' =
## 1.8651, : some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('7.692308%' = 1.57254615384615, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('7.692308%' = 1.57254615384615, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('7.692308%' = 1.57254615384615, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('7.692308%' = 1.57254615384615, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('7.142857%' = 1.53135, '14.28571%' =
## 1.7821, : some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('7.142857%' = 1.53135, '14.28571%' =
## 1.7821, : some 'x' values beyond boundary knots may cause ill-conditioned bases

```



```
## Warning in bs(dis, degree = 3L, knots = c('7.142857%' = 1.54498571428571, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('7.142857%' = 1.54498571428571, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('6.666667%' = 1.52093333333333, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('6.666667%' = 1.52093333333333, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases


## Warning in bs(dis, degree = 3L, knots = c('6.666667%' = 1.5218, '13.33333%' =
## 1.75478, : some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('6.666667%' = 1.5218, '13.33333%' =
## 1.75478, : some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('6.25%' = 1.5187, '12.5%' = 1.74615, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in bs(dis, degree = 3L, knots = c('6.25%' = 1.5187, '12.5%' = 1.74615, :
## some 'x' values beyond boundary knots may cause ill-conditioned bases

plot(3:18, deltas, type = 'l')
points(which.min(deltas) + 2, deltas[which.min(deltas)], col = 'red')
```



## ANOVA

```
fit1 = lm(nox ~ bs(dis, 3), data = Boston)
fit2 = lm(nox ~ bs(dis, 4), data = Boston)
fit3 = lm(nox ~ bs(dis, 5), data = Boston)
fit4 = lm(nox ~ bs(dis, 6), data = Boston)
fit5 = lm(nox ~ bs(dis, 7), data = Boston)
fit6 = lm(nox ~ bs(dis, 8), data = Boston)
fit7 = lm(nox ~ bs(dis, 9), data = Boston)
fit8 = lm(nox ~ bs(dis, 10), data = Boston)
fit9 = lm(nox ~ bs(dis, 11), data = Boston)
fit10 = lm(nox ~ bs(dis, 12), data = Boston)
fit11 = lm(nox ~ bs(dis, 13), data = Boston)
fit12 = lm(nox ~ bs(dis, 14), data = Boston)
fit13 = lm(nox ~ bs(dis, 15), data = Boston)
fit14 = lm(nox ~ bs(dis, 16), data = Boston)
fit15 = lm(nox ~ bs(dis, 17), data = Boston)
fit16 = lm(nox ~ bs(dis, 18), data = Boston)

anova(fit1, fit2, fit3, fit4, fit5, fit6, fit7, fit8, fit9,
      fit10, fit11, fit12, fit13, fit14, fit15, fit16)
```

```
## Analysis of Variance Table
##
```

```
## Model 1: nox ~ bs(dis, 3)
## Model 2: nox ~ bs(dis, 4)
## Model 3: nox ~ bs(dis, 5)
## Model 4: nox ~ bs(dis, 6)
## Model 5: nox ~ bs(dis, 7)
## Model 6: nox ~ bs(dis, 8)
## Model 7: nox ~ bs(dis, 9)
## Model 8: nox ~ bs(dis, 10)
## Model 9: nox ~ bs(dis, 11)
## Model 10: nox ~ bs(dis, 12)
## Model 11: nox ~ bs(dis, 13)
## Model 12: nox ~ bs(dis, 14)
## Model 13: nox ~ bs(dis, 15)
## Model 14: nox ~ bs(dis, 16)
## Model 15: nox ~ bs(dis, 17)
## Model 16: nox ~ bs(dis, 18)
##      Res.Df    RSS Df Sum of Sq      F    Pr(>F)
## 1      502 1.9341
## 2      501 1.9228  1  0.011332  3.1076 0.078556 .
## 3      500 1.8402  1  0.082602 22.6525 2.563e-06 ***
## 4      499 1.8340  1  0.006207  1.7022 0.192622
## 5      498 1.8299  1  0.004081  1.1193 0.290597
## 6      497 1.8170  1  0.012889  3.5347 0.060692 .
## 7      496 1.8256  1 -0.008657
## 8      495 1.7925  1  0.033118  9.0821 0.002716 **
## 9      494 1.7970  1 -0.004457
## 10     493 1.7890  1  0.007993  2.1919 0.139386
## 11     492 1.7824  1  0.006649  1.8233 0.177546
## 12     491 1.7818  1  0.000512  0.1405 0.707937
```

```
## 13      490 1.7828  1 -0.000960
## 14      489 1.7835  1 -0.000748
## 15      488 1.7798  1  0.003757  1.0303  0.310585
## 16      487 1.7758  1  0.003950  1.0833  0.298478
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

In general, using cross-validation or ANOVAm the results from polynomial regression and spline regression seem to agree with each other. An optimal degree varies from 4 to 6 is likely to yield a parsimonious model.

## Exercise 10

```
head(College)
```

```
##                               Private Apps Accept Enroll Top10perc Top25perc
## Abilene Christian University   Yes 1660  1232   721      23      52
## Adelphi University            Yes 2186  1924   512      16      29
## Adrian College                Yes 1428  1097   336      22      50
## Agnes Scott College           Yes  417   349   137      60      89
## Alaska Pacific University      Yes  193   146    55      16      44
## Albertson College             Yes  587   479   158      38      62
##                               F.Undergrad P.Undergrad Outstate Room.Board Books
## Abilene Christian University    2885           537   7440      3300  450
## Adelphi University              2683           1227  12280      6450  750
## Adrian College                 1036            99  11250      3750  400
## Agnes Scott College              510            63  12960      5450  450
## Alaska Pacific University        249           869   7560      4120  800
## Albertson College               678            41  13500      3335  500
##                               Personal PhD Terminal S.F.Ratio perc.alumni Expend
## Abilene Christian University    2200  70      78    18.1      12  7041
## Adelphi University             1500  29      30    12.2      16 10527
## Adrian College                 1165  53      66    12.9      30  8735
## Agnes Scott College              875  92      97     7.7      37 19016
## Alaska Pacific University       1500  76      72    11.9       2 10922
## Albertson College               675  67      73     9.4      11  9727
##                               Grad.Rate
## Abilene Christian University     60
## Adelphi University               56
## Adrian College                   54
## Agnes Scott College              59
## Alaska Pacific University         15
## Albertson College                 55
```

```
dim(na.omit(College))
```

```
## [1] 777  18
```

**a**

```

set.seed(1)
train_indices = sample(777, 444)
train_set = College[train_indices, ]
test_set = College[-train_indices, ]

```

```

reg_model = regsubsets(Outstate ~ ., data = train_set,
                       nvmax = 17, method = 'forward')
reg_summary = summary(reg_model)

```

```

par(mfrow = c(2, 2))

```

```

plot(reg_summary$rss, xlab = 'Number of Variables', ylab = 'RSS', type = 'l')
rss_min = which.min(reg_summary$rss)
points(rss_min, reg_summary$rss[rss_min], col = 'red', cex = 2, pch = 20)

```

```

plot(reg_summary$adjr2, xlab = 'Number of Variables', ylab = 'Adjusted Rsq', type = 'l')
adjr2_max = which.max(reg_summary$adjr2)
points(adjr2_max, reg_summary$adjr2[adjr2_max], col = 'red', cex = 2, pch = 20)

```

```

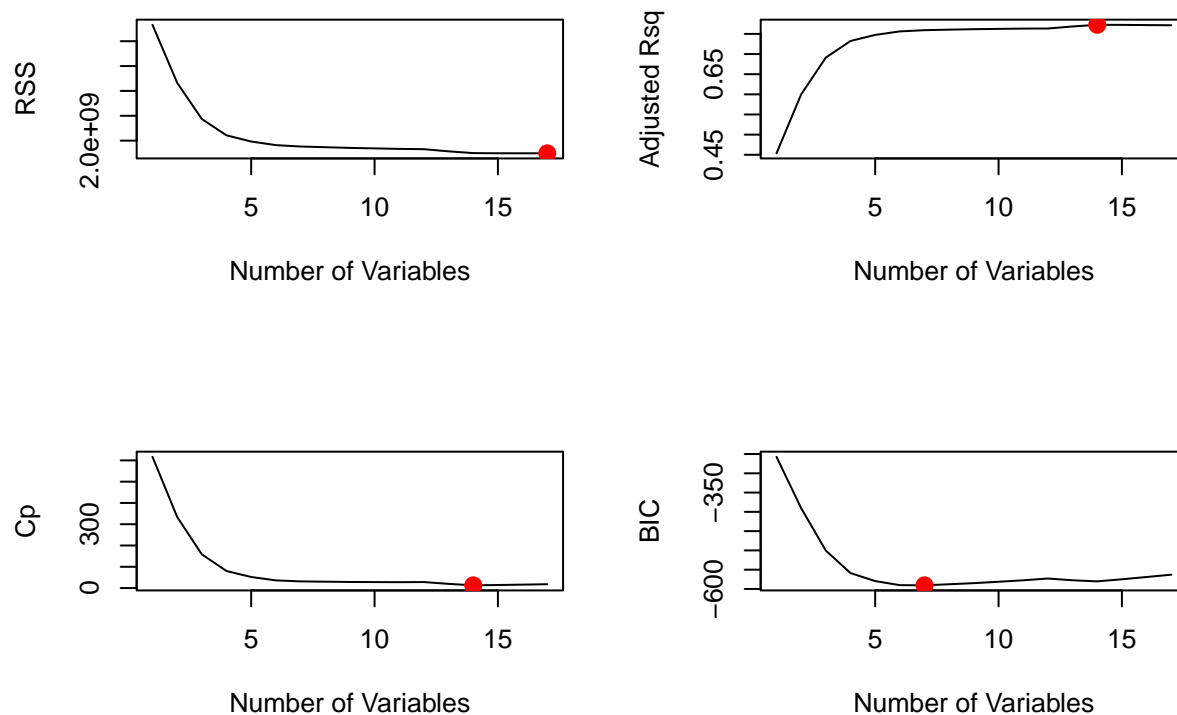
plot(reg_summary$cp, xlab = 'Number of Variables', ylab = 'Cp', type = 'l')
cp_min = which.min(reg_summary$cp)
points(cp_min, reg_summary$cp[cp_min], col = 'red', cex = 2, pch = 20)

```

```

plot(reg_summary$bic, xlab = 'Number of Variables', ylab = 'BIC', type = 'l')
bic_min = which.min(reg_summary$bic)
points(bic_min, reg_summary$bic[bic_min], col = 'red', cex = 2, pch = 20)

```



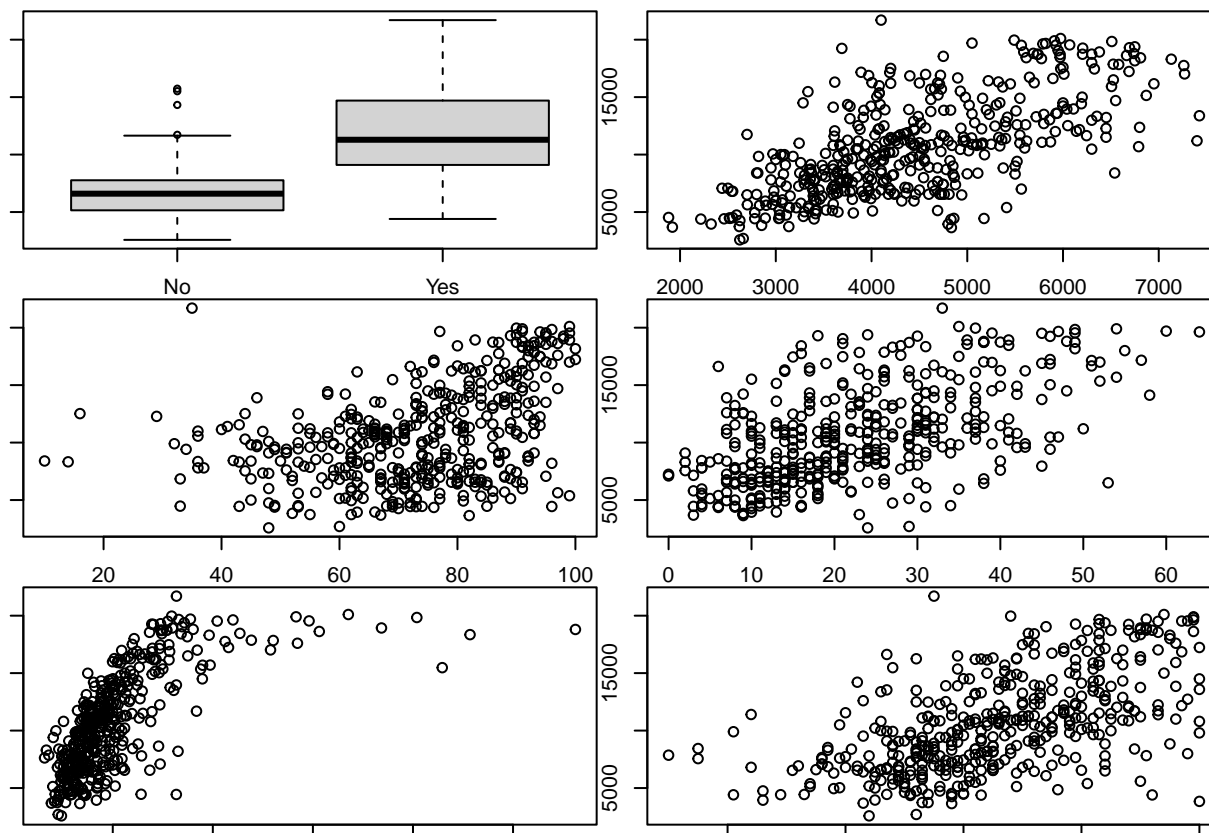
Here we would choose 6 as the optimal number of variables. It is understandable that using 15 or even all 17 variables results in small RSS or Adjusted R-squared since the model overfits the data. As you might notice, there is no significant difference in using from 6 to 17 variables. We want a parsimonious model.

```
coef(reg_model, 6)
```

```
##      (Intercept)    PrivateYes    Room.Board          PhD    perc.alumni
## -4065.9665184    2788.5353954      1.0799406      35.5263418      57.7080695
##      Expend      Grad.Rate
##      0.2013602      29.5272517
```


b

```
par(mfrow = c(3, 2))
par(mar = c(1, 1, 1, 1))
plot(train_set$Private, train_set$Outstate)
plot(train_set$Room.Board, train_set$Outstate)
plot(train_set$PhD, train_set$Outstate)
plot(train_set$perc.alumni, train_set$Outstate)
plot(train_set$Expend, train_set$Outstate)
plot(train_set$Grad.Rate, train_set$Outstate)
```



```
gam_model = gam(Outstate ~ Private + s(Room.Board, 5) + s(PhD, 5)
                + s(perc.alumni, 5) + s(Expend, 5) + s(Grad.Rate, 5),
                data = College)

par(mfrow = c(3, 2))
par(mar = c(1, 1, 1, 1))
plot(gam_model, se = T, col = 'blue')
```



c

```
linear_model = gam(Outstate ~ Private + Room.Board + PhD + perc.alumni
                    + Expend + Grad.Rate, data = College)
linear_preds = predict(linear_model, newdata = test_set)

err = mean((linear_preds - test_set$Outstate)^2)
tss = mean((mean(test_set$Outstate) - test_set$Outstate)^2)
1 - err / tss
```

```
## [1] 0.7268812
```

```
err
```

```
## [1] 3761153
```

```
gam_preds = predict(gam_model, newdata = test_set)
err = mean((gam_preds - test_set$Outstate)^2)
1 - err / tss
```

```
## [1] 0.7753438
```

```
err
```

```
## [1] 3093770
```

**d**

The plots and results from (c) suggest that there is a non-linear relationship between some predictors and the response. Specifically, they are Room.Board, PhD, Expend and Grad.Rate

```
summary(gam_model)
```

```
##
## Call: gam(formula = Outstate ~ Private + s(Room.Board, 5) + s(PhD,
##      5) + s(perc.alumni, 5) + s(Expend, 5) + s(Grad.Rate, 5),
##      data = College)
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -7576.49 -1102.48   47.67  1287.40  7492.43
##
## (Dispersion Parameter for gaussian family taken to be 3425681)
##
##      Null Deviance: 12559297426 on 776 degrees of freedom
## Residual Deviance: 2569258985 on 749.9994 degrees of freedom
## AIC: 13924.92
##
## Number of Local Scoring Iterations: NA
##
## Anova for Parametric Effects
##              Df      Sum Sq    Mean Sq F value    Pr(>F)
## Private              1 3370516566 3370516566 983.897 < 2.2e-16 ***
## s(Room.Board, 5)      1 2484278051 2484278051 725.192 < 2.2e-16 ***
## s(PhD, 5)             1  818616768  818616768 238.965 < 2.2e-16 ***
## s(perc.alumni, 5)     1 494085299 494085299 144.230 < 2.2e-16 ***
## s(Expend, 5)          1 1015946099 1015946099 296.568 < 2.2e-16 ***
## s(Grad.Rate, 5)       1 148600665 148600665  43.378 8.486e-11 ***
## Residuals           750 2569258985    3425681
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Anova for Nonparametric Effects
##              Npar Df Npar F    Pr(F)
## (Intercept)
## Private
## s(Room.Board, 5)      4  2.547 0.03824 *
## s(PhD, 5)             4  2.083 0.08126 .
## s(perc.alumni, 5)     4  1.144 0.33444
## s(Expend, 5)          4 32.545 < 2e-16 ***
## s(Grad.Rate, 5)       4  2.670 0.03121 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```





## Exercise 11

a

```
set.seed(1)

X1 = rnorm(100)
X2 = rnorm(100)
epsilon = rnorm(100)
Y = 1*X1 + 2*X2 + 3 + epsilon
```

```
par(mfrow = c(1, 2))
plot(X1, Y)
plot(X2, Y)
```



b

```
B1_hat = 100
```

**c**

```
a = Y - B1_hat * X1
B2_hat = lm(a ~ X2)$coef[2]
B2_hat
```

```
##          X2
## 2.038818
```

**d**

```
a = Y - B2_hat * X2
B1_hat = lm(a ~ X1)$coef[2]
B1_hat
```

```
##          X1
## 1.021208
```

**e**

Here I'll only iterate the process for 10 times since we got very close to the actual coefficients just after 2 tries.

```
B1_hat = 10

B0s_hat = rep(NA, 10)
B1s_hat = rep(NA, 10)
B2s_hat = rep(NA, 10)

for (i in 1:10){


  a = Y - B1_hat * X1
  B2_hat = lm(a ~ X2)$coef[2]
  B2s_hat[i] = B2_hat

  a = Y - B2_hat * X2
  model = lm(a ~ X1)
  B1_hat = model$coef[2]
  B1s_hat[i] = B1_hat

  B0s_hat[i] = model$coef[1]

}
```

```
par(mfrow = c(1, 3))
plot(1:10, B0s_hat)
plot(1:10, B1s_hat)
plot(1:10, B2s_hat)
```




f

```
multi_model = lm(Y ~ X1 + X2)
coeff = multi_model$coef
coeff
```

```
## (Intercept)      X1      X2
##    3.025353    1.021110    1.946533
```

```
par(mfrow = c(1, 3))

plot(1:10, B0s_hat, col = 'blue')
abline(h = coeff[1], col = 'green')
plot(1:10, B1s_hat, col = 'blue')
abline(h = coeff[2], col = 'green')
plot(1:10, B2s_hat, col = 'blue')
abline(h = coeff[3], col = 'green')
```



g

We only need a single backfitting iteration to obtain a “good” approximation to the multiple regression coefficient estimates.

## Exercise 12

```
set.seed(1)

B = c(1:100)
intercept = 1
X = matrix(rnorm(1000 * 100), nrow = 1000, ncol = 100)
e = rnorm(1000)

Y = intercept + X %*% B + e
```

```
set.seed(1)

B_hat = rnorm(100)
iters = 30
B_his = matrix(NA, nrow = iters, ncol = 101)
```

```

for (i in 1:iters) {

  for (p in 1:100) {

    a = Y - X[, -p] %*% B_hat[-p]
    coeff = lm(a ~ X[, p])$coef
    B_hat[p] = coeff[2]

    B_his[i, p+1] = B_hat[p]
  }

  beta0 = coeff[1]
  B_his[i, 1] = beta0
}

```

```
B_hat[1:5]
```

```
## [1] 0.9727505 2.0606211 2.9936625 3.9582157 4.9714572
```

```
B[1:5]
```

```
## [1] 1 2 3 4 5
```

```
B_his[1:20, 1:6]
```

```

##           [,1]      [,2]      [,3]      [,4]      [,5]      [,6]
## [1,] -0.9928798 -2.1347165 1.080987 1.800493 -11.531064 -29.444173
## [2,] 0.7535299 5.2524427 4.971155 -5.542070 6.152879 1.103353
## [3,] 0.8613899 2.0157122 2.894328 2.137096 2.890126 3.784716
## [4,] 0.9998247 1.0441617 1.934883 3.017297 3.712200 4.802683
## [5,] 1.0189743 0.9563286 1.975577 3.061277 3.911303 4.935899
## [6,] 1.0201238 0.9620477 2.021010 3.024469 3.945772 4.965000
## [7,] 1.0191109 0.9694400 2.046744 3.003595 3.954409 4.970677
## [8,] 1.0185902 0.9718723 2.056407 2.996406 3.957068 4.971361
## [9,] 1.0184064 0.9725178 2.059460 2.994357 3.957918 4.971422
## [10,] 1.0183484 0.9726867 2.060324 2.993826 3.958154 4.971439
## [11,] 1.0183311 0.9727330 2.060550 2.993698 3.958208 4.971449
## [12,] 1.0183262 0.9727459 2.060605 2.993669 3.958217 4.971454
## [13,] 1.0183249 0.9727494 2.060618 2.993663 3.958217 4.971456
## [14,] 1.0183245 0.9727502 2.060621 2.993662 3.958216 4.971457
## [15,] 1.0183244 0.9727504 2.060621 2.993662 3.958216 4.971457
## [16,] 1.0183244 0.9727505 2.060621 2.993662 3.958216 4.971457
## [17,] 1.0183244 0.9727505 2.060621 2.993662 3.958216 4.971457
## [18,] 1.0183244 0.9727505 2.060621 2.993663 3.958216 4.971457
## [19,] 1.0183244 0.9727505 2.060621 2.993663 3.958216 4.971457
## [20,] 1.0183244 0.9727505 2.060621 2.993663 3.958216 4.971457

```

The algorithm obtained “good” approximation to coefficients just after around 5 iterations. Cool :)